

# Ten simple rules to follow when cleaning occurrence data in palaeobiology

The Palaeoverse Development Team

## Contents

Introduction . . . . .	1
Load packages . . . . .	1
Rule 1: Choose the right data for your question . . . . .	2
Rule 2: Keep raw data raw . . . . .	2
Rule 3: Document your workflow . . . . .	4
Rule 4: Explore your data . . . . .	4
Rule 5: Identify and handle incomplete data records . . . . .	6
Rule 6: Identify and handle outliers . . . . .	7
Rule 7: Identify and handle inconsistencies . . . . .	13
Rule 8: Identify and handle duplicates . . . . .	27
Rule 9: Report your data and cleaning efforts . . . . .	29
Rule 10: Deposit your data and workflow . . . . .	30
Summary . . . . .	30

## Introduction

This vignette is an accompaniment to:

Jones LA, Dean CD, Allen BJ, Drage HB, Flannery-Sutherland JT, Gearty W, Chiarenza AA, Dillon EM, Farina BM, and Godoy PL. Ten simple rules to follow when cleaning occurrence data in palaeobiology.

Here we present a full example workflow that complies with the ten rules we propose in the manuscript. The intention is to illustrate the concepts outlined in the manuscript, with further details and ideas for practical implementation, as well as code that can be repurposed by readers.

## Load packages

Before starting, we will load all of the R packages we need.

```
# install.packages(c("CoordinateCleaner", "deeptime", "dplyr", "fossilbrush",  
# "ggplot2", "palaeoverse", "readr", "rgplates", "rnaturalearth",  
# "rnaturalearthdata"))  
library(CoordinateCleaner)  
library(deeptime)  
library(dplyr)  
library(fossilbrush)  
library(ggplot2)  
library(palaeoverse)  
library(readr)  
library(rgplates)
```

```
library(rnaturalearth)
library(rnaturalearthdata)
```

## Rule 1: Choose the right data for your question

In this example, we are interested in the fossil record of crocodiles. As ectotherms, crocodiles are highly reliant on the environment in which they live in order to maintain a functional internal body temperature. Because of this, their spatial distribution is constrained to warm climates, so their fossil record is commonly used as an indicator of palaeoclimate. Here we will investigate the palaeodiversity of crocodiles, with the specific goal of reconstructing their latitudinal range throughout the Paleogene.

To meet this goal, we need to acquire occurrence data for fossil crocodiles during the Paleogene. Initially we have decided not to place further taxonomic constraints on our search, so we will include all occurrences belonging to the order ‘Crocodylia’. We are interested in the clade’s biogeography, so we will need all occurrences globally, and we need to ensure that we have geographic coordinates associated with our occurrences.

We will turn to one of the largest sources of fossil occurrence data, the [Paleobiology Database](#). We opted to download our data by completing the ‘Download’ form on the Paleobiology Database website. Practically, this uses the options entered into the form to create a call to the database’s API (application programming interface), and then enacts this call. More information about the API service can be found [here](#).

For our dataset, we pulled all occurrences associated with the taxon name ‘Crocodylia’, dated to the ‘Paleogene’. All other settings were left at default. We obtained the standard (or recommended) data fields by checking the box to “Include all output blocks whose names are boldfaced below”, and obtained our data in the form of a ‘csv’ file. We will see later in the vignette what this means in terms of what is contained within the dataset.

## Rule 2: Keep raw data raw

For reproducibility, we want to make sure that we have a copy of the full dataset as initially downloaded - this is the “raw” data. Automatically, when we use the online form, there is an output option checked which is labelled “Include metadata at the beginning of the output”. This ensures that the raw data file includes a metadata ‘header’, which contains useful information linked to the dataset.

We can load and view the metadata to see what it includes.

```
# Load data file
metadata <- read_csv("data/Paleogene_crocs.csv")
```

```
# Trim to metadata
metadata <- metadata[1:23, ]
```

```
# Print
metadata
```

```
## # A tibble: 23 x 2
##   'Data Provider'      'The Paleobiology Database'
##   <chr>                <chr>
## 1 Data Source          The Paleobiology Database
## 2 Data License          Creative Commons CC0
## 3 License URL          https://creativecommons.org/publicdomain/zero/1.0/
## 4 Documentation URL    http://paleobiodb.org/data1.2/occs/list_doc.html
## 5 Data URL             http://paleobiodb.org/data1.2/occs/list.csv?datainfo&rowco~
## 6 Access Time          Fri 2025-02-21 14:41:31 GMT
## 7 Title                PBDB Data Service
```

```
## 8 Parameters:      <NA>
## 9 <NA>             base_name,Crocodylia
## 10 <NA>            pgm,gplates,scotese,seton
## # i 13 more rows
```

The metadata are strangely formatted here, but we can see that they include information about the data license (CC0), the API call used (under the label ‘Data URL’), the date and time at which the data were accessed, and the total number of records contained within the dataset (here, 886 fossil occurrences).

These metadata elements are all important information to retain alongside our data, allowing others to better understand what the dataset contains, and when and how it was downloaded. The Paleobiology Database is fully dynamic, not only in that new data is continually being added, but also in that any record can be changed retrospectively by an Editor. It cannot be assumed that the ‘present’ state of any data record was the same in the (historical) past. So, for example, if someone wanted to see how the data associated with this API call had changed in the time elapsed since our download, they could do this, and directly in R if desired:

```
# View API URL
metadata[5, 2]

## # A tibble: 1 x 1
##   ‘The Paleobiology Database‘
##   <chr>
## 1 http://paleobiodb.org/data1.2/occs/list.csv?datainfo&rowcount&base_name=Croco~

# Use API call (this is not enacted here)
new_data <- read_csv(metadata[5, 2])
```

While the metadata is important to keep in the raw file, for the purposes of analysis, we want to be able to just read in the data beneath it. We can do this using the `skip` parameter in `read_csv`, which tells R to ignore a given number of rows at the top of the file.

```
# Load data file, skipping metadata
fossils <- read_csv("data/Paleogene_crocs.csv", skip = 19)

## New names:
## Rows: 886 Columns: 131
## -- Column specification
## ----- Delimiter: "," chr
## (95): record_type, flags, identified_name, identified_rank, difference, ... dbl
## (20): occurrence_no, reid_no, collection_no, identified_no, accepted_no,... lgl
## (16): accepted_attr, plant_organ, plant_organ2, regionalbedunit, artifa...
## i Use ‘spec()’ to retrieve the full column specification for this data. i
## Specify the column types or set ‘show_col_types = FALSE’ to quiet this message.
## * ‘cc’ -> ‘cc...36’
## * ‘cc’ -> ‘cc...58’
```

When we use `read_csv()`, we get a message explaining how the data have been parsed into R. It’s worth checking this for anything unusual, because if parsing has not occurred how we expected, it could lead to errors in the data. Here, we can see that there were two columns in the `csv` file named “cc”, for “country code”. Their column number has been appended to their column name, in order to keep these distinct. Is this column simply duplicated? We can check this.

```
# Are the two `cc` columns identical?
identical(fossils$cc...36, fossils$cc...58)
```

```
## [1] TRUE
```

This is true, so to keep our dataframe tidy, we will remove one of these columns and rename the other.

```
# Remove one `cc` column
fossils$cc...58 <- NULL

# Rename other column
colnames(fossils)[colnames(fossils) == 'cc...36'] <- 'cc'
```

And now we are ready to commence our data exploration and cleaning.

### Rule 3: Document your workflow

Documenting your workflow is essential for ensuring that others can understand and replicate all steps. Using programming languages such as R or Python makes this easy, as the code used can be provided alongside the data. Our vignette gives an example of how **markdown** can be a particularly useful tool for integrating code with prose, ensuring that everything is fully explained, and providing plots integrated within the text when exported.

Our tips for good documentation in R include:

- Following general guidelines for good coding, such as ample use of spacing to help human readability
- Having a clear, overarching code structure (such as subsections, which can be denoted with headings and “#====” notation, and an order which corresponds to the flow of the research methods/results)
- Using specific yet short object names
- Including comments throughout which explain both *what* and *why*

### Rule 4: Explore your data

The first thing we want to do with our data is *explore* it - that is, generate summary statistics and plots to help us understand the data and its various characteristics.

For example, we can look at the distribution of identification levels for our fossils.

```
# Count the frequency of taxonomic ranks
table(fossils$accepted_rank)
```

```
##
##      family      genus      species      subfamily      superfamily
##          1          254          250           3           13
## unranked clade
##          365
```

```
# Calculate as percentages
(table(fossils$accepted_rank) / nrow(fossils)) * 100
```

```
##
##      family      genus      species      subfamily      superfamily
## 0.1128668 28.6681716 28.2167043 0.3386005 1.4672686
## unranked clade
## 41.1963883
```

We can see that of our 886 occurrences, 250 (28%) are identified to species level. A further 254 (29%) are identified to genus level. The remaining fossils are more coarsely identified, including 365 (41%) which are identified to the mysterious level of “unranked clade”.

Next, let’s look at the distribution of fossils across localities. In the PBDB, fossils are placed within collections, each of which can roughly be considered a separate locality (they can also represent different sampling horizons at the same locality; more on this later). First, we can count the number of unique `collection_no` values to find out how many unique collections are in the dataset.


```
# What is the length of a vector of unique collection numbers?
length(unique(fossils$collection_no))
```

```
## [1] 720
```

Our dataset contains 720 unique collections. We can also create a plot showing us the distribution of occurrences across these collections.

```
# Count the number of times each collection number appears in the dataset
coll_no_freq <- as.data.frame(table(fossils$collection_no))

# Plot the distribution of number of occurrences per collection
ggplot(coll_no_freq, aes(x = Freq)) +
  geom_bar() +
  labs(x = "Number of occurrences",
       y = "Frequency")
```



We can see that the collection containing the most occurrences has 7, while the vast majority only contain a single occurrence.

What about the countries in which these fossils were found? We can investigate this using the “cc”, or “country code” column.

```
# List unique country codes, and count them
unique(fossils$cc)
```

```
## [1] "US" "CA" "PE" "IN" "FR" "MA" "BE" "RU" "AR" "ES" "UK" "DE" "KZ" "LY" "IT"
## [16] "BR" "AU" "PK" "JM" "EG" "AA" "TR" "CZ" "MG" "CN" "DK" "SE" "UA" "PL" "PR"
## [31] "TH" "TN" "VN" "KE" "AT" "OM" "PT" NA "CH" "EH" "ML" "RO" "NP" "MN" "BG"
## [46] "CO"
```

```
length(unique(fossils$cc))
```

```
## [1] 46
```

Here we can see that Paleogene crocodiles have been found in 46 different countries. Let's sort those values alphabetically to help us find specific countries.

```
# List and sort unique country codes, and count them  
sort(unique(fossils$cc))
```

```
## [1] "AA" "AR" "AT" "AU" "BE" "BG" "BR" "CA" "CH" "CN" "CO" "CZ" "DE" "DK" "EG"  
## [16] "EH" "ES" "FR" "IN" "IT" "JM" "KE" "KZ" "LY" "MA" "MG" "ML" "MN" "NP" "OM"  
## [31] "PE" "PK" "PL" "PR" "PT" "RO" "RU" "SE" "TH" "TN" "TR" "UA" "UK" "US" "VN"
```

```
length(sort(unique(fossils$cc)))
```

```
## [1] 45
```

Something weird has happened here: we can see that once the countries have been sorted, one of them has disappeared. Why? We will come back to this during our next rule.

## Rule 5: Identify and handle incomplete data records

By default, when we read data tables into R, it recognises empty cells and takes some course of action to manage them. When we use base R functions, such as `read.csv()`, empty cells are given an NA value ('not available') only when the column is considered to contain numerical data. When we use **Tidyverse** functions, such as `readr::read_csv()`, *all* empty cells are given NA values. This is important to bear in mind when we want to find those missing values: here, we have done the latter, so all empty cells are NA.

The extent of incompleteness of the different columns in our dataset is highly variable. For example, the number of NA values for the `collection_no` is 0.

```
# Count the number of collection number values for which `is.na()` is TRUE  
sum(is.na(fossils$collection_no))
```

```
## [1] 0
```

This is because it is impossible to add an occurrence to the PBDB without putting it in a collection, which must in turn have an identification number.

However, what about `genus`?

```
# Count the number of genus IDs for which `is.na()` is TRUE  
sum(is.na(fossils$genus))
```

```
## [1] 382
```

The value here is 382. This corresponds to the number of occurrences in our dataset which cannot be placed in a genus, i.e. their taxonomic identification is to a coarser taxonomic level.

What can, or should, we do about this? It depends heavily on the intersection between the different fields in the dataset, and our research question. In some instances it might be appropriate to ignore certain NA values: for example, NA values for `genus` might not matter if we want to see the spatial distribution of all crocodylian fossils, regardless of how coarsely identified they are.

Alternatively, it might be appropriate to filter these values out. For our crocodile example, we are interested in biogeography, so the latitude and longitude, and palaeolatitude and palaeolongitude, are pretty important. Do we have missing values there?

```
# Count the number of occurrences for which `is.na()` is TRUE in latitude,  
# longitude, palaeolatitude, and palaeolongitude  
sum(is.na(fossils$lat))
```

```
## [1] 0
```

```
sum(is.na(fossils$lng))
```

```
## [1] 0
```

```
sum(is.na(fossils$paleolat))
```

```
## [1] 18
```

```
sum(is.na(fossils$paleolng))
```

```
## [1] 18
```

While all occurrences have modern day coordinates, 18 are missing palaeocoordinates. We will now remove these 18 occurrences from the dataset.

```
# Remove occurrences which are missing palaeocoordinates  
fossils <- filter(fossils, !is.na(fossils$paleolng))
```

```
# Check whether this has worked  
sum(is.na(fossils$paleolat))
```

```
## [1] 0
```

```
sum(is.na(fossils$paleolng))
```

```
## [1] 0
```

A further option applicable in some cases would be to fill in our missing data. We may be able to interpolate values from the rest of our data, or use additional data sources. For our palaeogeography example above, we could generate our own coordinates, for example using `palaeoverse::palaeorotate()`.

Let's revisit our country example from Rule 5. Does our list of country codes contain NAs?

```
# Count the number of country codes for which `is.na()` is TRUE  
sum(is.na(fossils$cc))
```

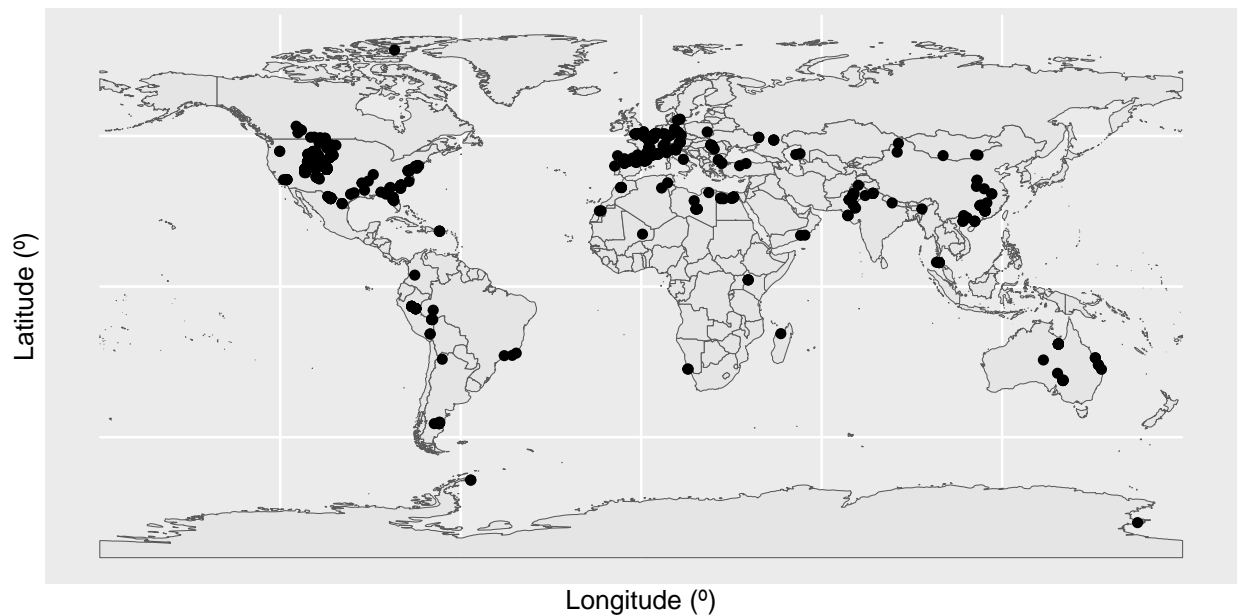
```
## [1] 2
```

Our test tells us that two of the values are NA: when we asked R to sort the values, it removed NA from the list of unique countries. However, the PBDB has set the country within which the collection is located as a compulsory entry field. How can we have missing data here? We don't: these values are not NA, they are "NA" meaning Namibia, and have been misconstrued by R. This is an important illustration of why we should conduct further investigation when any apparent errors arise in the dataset, rather than immediately removing these data points.

## Rule 6: Identify and handle outliers

Next, we want to look for outliers. Here we will focus in on the specific variables which relate to our scientific question, i.e. the geography of our fossil occurrences. First we'll plot where the crocodile fossils have been found across the globe: how does this match what we already know from the country codes?

```
# Load in a world map  
world <- ne_countries(scale = "medium", returnclass = "sf")  
  
# Plot the geographic coordinates of each locality over the world map  
ggplot(fossils) +  
  geom_sf(data = world) +  
  geom_point(aes(x = lng, y = lat)) +  
  labs(x = "Longitude (°)",  
       y = "Latitude (°)")
```




We have a large density of crocodile occurrences in Europe and the western interior of the United States, along with a smattering of occurrences across the other continents. This distribution seems to fit our previous knowledge, that the occurrences are spread across 46 countries. However, the crocodile occurrences in Antarctica seem particularly suspicious: crocodiles need a warm climate, and modern-day Antarctica certainly doesn't fit this description. Let's investigate further. We'll do this by plotting the latitude of the occurrences through time.

```
# Add a column to the data frame with the midpoint of the fossil ages
fossils <- mutate(fossils, mid_ma = (min_ma + max_ma) / 2)

# Create dataset containing only Antarctic fossils
Antarctic <- filter(fossils, cc == "AA")

# Plot the age of each occurrence against its latitude
ggplot(fossils, aes(x = mid_ma, y = lat)) +
  geom_point(colour = "black") +
  geom_point(data = Antarctic, colour = "red") +
  labs(x = "Age (Ma)",
       y = "Latitude (°)") +
  scale_x_reverse() +
  geom_hline(yintercept = 0) +
  coord_geo(dat = "stages", expand = TRUE)
```






Here we can see the latitude of each occurrence, plotted against the temporal midpoint of the collection. We have highlighted our Antarctic occurrences in red - these points are still looking pretty anomalous.

But, wait, we should actually be looking at **palaeolatitude** instead. Let's plot that against time.

```
# Plot the age of each occurrence against its palaeolatitude
ggplot(fossils, aes(x = mid_ma, y = paleolat)) +
  geom_point(colour = "black") +
  geom_point(data = Antarctic, colour = "red") +
  labs(x = "Age (Ma)",
       y = "Palaeolatitude (°)") +
  scale_x_reverse() +
  geom_hline(yintercept = 0) +
  coord_geo(dat = "stages", expand = TRUE)
```



Hmm... when we look at palaeolatitude the Antarctic occurrences are even **further** south. Time to really check out these occurrences. Which collections are they within?

```
# Find Antarctic collection numbers
unique(Antarctic$collection_no)
```

```
## [1] 43030 120887 31173
```

Well, upon further visual inspection using the PBDB website, all appear to be fairly legitimate. However, all three occurrences still appear to be outliers, especially as in the late Eocene [temperatures were dropping](#). What about the taxonomic certainty of these occurrences?

```
# List taxonomic names associated with Antarctic occurrences
Antarctic$identified_name
```

```
## [1] "Crocodylia indet." "Crocodylia indet." "Crocodylia indet."
```

Since all three occurrences are listed as “Crocodylia indet.”, it may make sense to remove them from further analyses anyway.


Let’s investigate if there are any other anomalies or outliers in our data. We’ll bin the occurrences by stage to look for stage-level outliers, using boxplots to show us any anomalous data points.

```
# Put occurrences into stage bins
bins <- time_bins(scale = "international ages")
fossils <- bin_time(occdf = fossils, bins = bins,
  min_ma = "min_ma", max_ma = "max_ma", method = "majority")

# Add interval name labels to occurrences
bins <- select(bins, bin, interval_name)
fossils <- left_join(fossils, bins, by = c("bin_assignment" = "bin"))

# Plot occurrences
ggplot(fossils, aes(x = bin_midpoint, y = paleolat, fill = interval_name)) +
  geom_boxplot(show.legend = FALSE) +
```

```
labs(x = "Age (Ma)",
     y = "Palaeolatitude (°)") +
scale_x_reverse() +
scale_fill_geo("stages") +
coord_geo(dat = "stages", expand = TRUE)
```

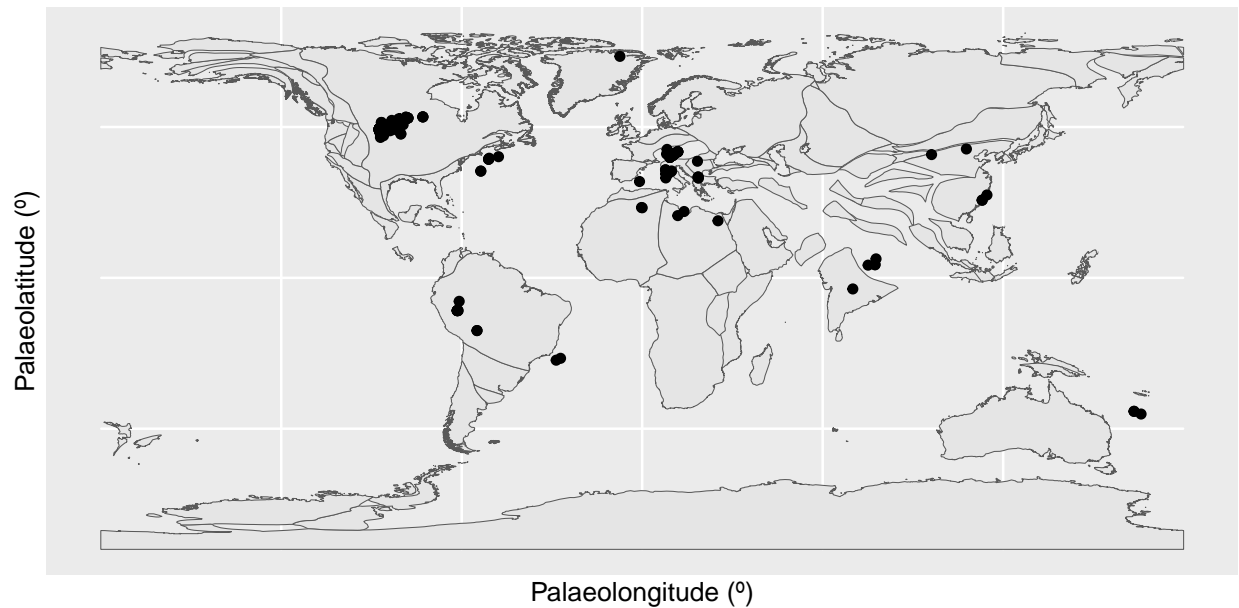


Box plots are a great way to look for outliers, because their calculation automatically includes outlier determination, and any such points can clearly be seen in the graph. At time of writing, the guidance for `geom_boxplot()` states that “The upper whisker extends from the hinge to the largest value no further than  $1.5 \times \text{IQR}$  from the hinge (where IQR is the inter-quartile range, or distance between the first and third quartiles). The lower whisker extends from the hinge to the smallest value at most  $1.5 \times \text{IQR}$  of the hinge. Data beyond the end of the whiskers are called ‘outlying’ points and are plotted individually.” 1.5 times the interquartile range seems a reasonable cut-off for determining outliers, so we will use these plots at face value to identify data points to check.

Here, the Ypresian (“Y”) is looking pretty suspicious - it seems to have a lot of outliers. Let’s plot the Ypresian occurrences on a palaeogeographic map to investigate further.

```
# Load map of the Ypresian, and identify Ypresian fossils
fossils_y <- fossils %>%
  filter(interval_name == "Ypresian")
world_y <- reconstruct("coastlines", model = "PALEOMAP", age = 51.9)


# Plot localities on the Ypresian map
ggplot(fossils_y) +
  geom_sf(data = world_y) +
  geom_point(aes(x = paleolong, y = paleolat)) +
  labs(x = "Palaeolongitude (°)",
       y = "Palaeolatitude (°)")
```



Aha! There is a concentrated cluster of occurrences in the western interior of North America. This high number of occurrences is increasing the weight of data at this palaeolatitude, and narrowing the boundaries at which other points are considered outliers. We can check the effect this is having on our outlier identification by removing the US occurrences from the dataset and checking the distribution again.

```
# Remove US fossils from the Ypresian dataset
fossils_y <- fossils_y %>%
  filter(cc != "US")

# Plot boxplot of non-US Ypresian fossil palaeolatitudes
ggplot(fossils_y) +
  geom_boxplot(aes(y = paleolat)) +
  labs(y = "Palaeolatitude (°)") +
  scale_x_continuous(breaks = NULL)
```



We can now see that none of our occurrences are being flagged as outliers. Without this strong geographic bias towards the US, all of the occurrences in the Ypresian appear to be reasonable. This fits our prior knowledge, as [elevated global temperatures during this time](#) likely helped crocodiles to live at higher latitudes than was possible earlier in the Paleogene.

So to sum up, it seems that our outliers are not concerning, so we will leave them in our dataset and continue with our analytical pipeline.

## Rule 7: Identify and handle inconsistencies

We're now going to look for inconsistencies in our dataset. Let's start by revisiting its structure, focusing on whether the class types of the variables make sense.

```
# Check the data class of each field in our dataset
str(fossils)
```

```
## tibble [868 x 137] (S3: tbl_df/tbl/data.frame)
##  $ occurrence_no      : num [1:868] 40163 40167 40168 40169 150323 ...
##  $ record_type        : chr [1:868] "occ" "occ" "occ" "occ" ...
##  $ reid_no            : num [1:868] 18506 NA NA NA NA ...
##  $ flags              : chr [1:868] NA NA NA NA ...
##  $ collection_no      : num [1:868] 3113 3113 3113 3113 13346 ...
##  $ identified_name     : chr [1:868] "Crocodylia indet." "Thoracosaurus basifissus" "Thoracosaurus l..."
##  $ identified_rank     : chr [1:868] "unranked clade" "species" "species" "species" ...
##  $ identified_no       : num [1:868] 38309 216615 216614 184628 38435 ...
##  $ difference         : chr [1:868] NA "nomen dubium" "nomen dubium" NA ...
##  $ accepted_name      : chr [1:868] "Crocodylia" "Gavialoidea" "Gavialoidea" "Thoracosaurus neoces..."
##  $ accepted_attr      : logi [1:868] NA NA NA NA NA NA ...
##  $ accepted_rank      : chr [1:868] "unranked clade" "superfamily" "superfamily" "species" ...
##  $ accepted_no        : num [1:868] 36582 96627 96627 184627 38435 ...
##  $ early_interval     : chr [1:868] "Thanetian" "Thanetian" "Thanetian" "Thanetian" ...
##  $ late_interval      : chr [1:868] NA NA NA NA ...
##  $ max_ma             : num [1:868] 59.2 59.2 59.2 59.2 47.8 54.9 66 66 63.8 54.9 ...
```

```

## $ min_ma : num [1:868] 56 56 56 56 41.2 50.5 61.6 61.6 60.9 50.5 ...
## $ ref_author : chr [1:868] "Alroy" "Cook and Ramsdell" "Cook and Ramsdell" "Cook and Ramsdell" ...
## $ ref_pubyr : num [1:868] 2006 1991 1991 1991 1988 ...
## $ reference_no : num [1:868] 18120 140 140 140 688 ...
## $ phylum : chr [1:868] "Chordata" "Chordata" "Chordata" "Chordata" ...
## $ class : chr [1:868] "Reptilia" "Reptilia" "Reptilia" "Reptilia" ...
## $ order : chr [1:868] "Crocodylia" "Crocodylia" "Crocodylia" "Crocodylia" ...
## $ family : chr [1:868] NA NA NA "Gavialidae" ...
## $ genus : chr [1:868] NA NA NA "Thoracosaurus" ...
## $ plant_organ : logi [1:868] NA NA NA NA NA NA ...
## $ plant_organ2 : logi [1:868] NA NA NA NA NA NA ...
## $ abund_value : num [1:868] NA NA NA NA 62 NA NA NA NA ...
## $ abund_unit : chr [1:868] NA NA NA NA ...
## $ lng : num [1:868] -74.7 -74.7 -74.7 -74.7 -86.5 ...
## $ lat : num [1:868] 40 40 40 40 31.4 ...
## $ occurrence_comments : chr [1:868] "originally entered as \"Crocodylus? sp.\" NA NA NA ...
## $ collection_name : chr [1:868] "Vincentown Formation, NJ" "Vincentown Formation, NJ" "Vincentown Formation, NJ" ...
## $ collection_subset : num [1:868] NA NA NA NA NA NA NA NA NA ...
## $ collection_aka : chr [1:868] NA NA NA NA ...
## $ cc : chr [1:868] "US" "US" "US" "US" ...
## $ state : chr [1:868] "New Jersey" "New Jersey" "New Jersey" "New Jersey" ...
## $ county : chr [1:868] NA NA NA NA ...
## $ latlng_basis : chr [1:868] "estimated from map" "estimated from map" "estimated from map" "estimated from map" ...
## $ latlng_precision : chr [1:868] "seconds" "seconds" "seconds" "seconds" ...
## $ geogscale : chr [1:868] "local area" "local area" "local area" "local area" ...
## $ geogcomments : chr [1:868] "\"The Vincentown Fm. occurs in an irregular, narrow belt extending from the base of the ...
## $ paleomodel : chr [1:868] "gplates" "gplates" "gplates" "gplates" ...
## $ geoplate : chr [1:868] "109" "109" "109" "109" ...
## $ paleoage : chr [1:868] "mid" "mid" "mid" "mid" ...
## $ paleolng : num [1:868] -44.5 -44.5 -44.5 -44.5 -66.8 ...
## $ paleolat : num [1:868] 40.1 40.1 40.1 40.1 34.7 ...
## $ paleomodel2 : chr [1:868] "scotese" "scotese" "scotese" "scotese" ...
## $ geoplate2 : num [1:868] 109 109 109 109 109 101 101 101 133 101 ...
## $ paleoage2 : chr [1:868] "mid" "mid" "mid" "mid" ...
## $ paleolng2 : num [1:868] -55.1 -55.1 -55.1 -55.1 -75.5 ...
## $ paleolat2 : num [1:868] 40.1 40.1 40.1 40.1 32.4 ...
## $ paleomodel3 : chr [1:868] "seton" "seton" "seton" "seton" ...
## $ geoplate3 : chr [1:868] "101" "101" "101" "101" ...
## $ paleoage3 : chr [1:868] "mid" "mid" "mid" "mid" ...
## $ paleolng3 : num [1:868] -54 -54 -54 -54 -75.3 ...
## $ paleolat3 : num [1:868] 39.3 39.3 39.3 39.3 29.7 ...
## $ protected : chr [1:868] NA NA NA NA ...
## $ formation : chr [1:868] "Vincentown" "Vincentown" "Vincentown" "Vincentown" ...
## $ stratgroup : chr [1:868] NA NA NA NA ...
## $ member : chr [1:868] NA NA NA NA ...
## $ stratscale : chr [1:868] "formation" "formation" "formation" "formation" ...
## $ zone : chr [1:868] NA NA NA NA ...
## $ localsection : chr [1:868] "New Jersey" "New Jersey" "New Jersey" "New Jersey" ...
## $ localbed : chr [1:868] NA NA NA NA ...
## $ localbedunit : chr [1:868] NA NA NA NA ...
## $ localorder : chr [1:868] NA NA NA NA ...
## $ regionalsection : chr [1:868] NA NA NA NA ...
## $ regionalbed : chr [1:868] NA NA NA NA ...
## $ regionalbedunit : logi [1:868] NA NA NA NA NA NA ...

```

```
## $ regionalorder      : chr [1:868] NA NA NA NA ...
## $ stratcomments      : chr [1:868] NA NA NA NA ...
## $ lithdescript       : chr [1:868] NA NA NA NA ...
## $ lithology1         : chr [1:868] "sandstone" "sandstone" "sandstone" "sandstone" ...
## $ lithadj1           : chr [1:868] "glaconitic" "glaconitic" "glaconitic" "glaconitic" ...
## $ lithification1     : chr [1:868] NA NA NA NA ...
## $ minor_lithology1   : chr [1:868] "sandy,calcareous" "sandy,calcareous" "sandy,calcareous" "sandy,calcareous" ...
## $ fossilsfrom1       : chr [1:868] NA NA NA NA ...
## $ lithology2         : chr [1:868] NA NA NA NA ...
## $ lithadj2           : chr [1:868] NA NA NA NA ...
## $ lithification2     : chr [1:868] NA NA NA NA ...
## $ minor_lithology2   : chr [1:868] NA NA NA NA ...
## $ fossilsfrom2       : chr [1:868] NA NA NA NA ...
## $ environment        : chr [1:868] NA NA NA NA ...
## $ tectonic_setting   : chr [1:868] NA NA NA NA ...
## $ geology_comments   : chr [1:868] "lithology described as a calcareous \"lime sand\" interbedded with a silty claystone"
## $ assembl_comps      : chr [1:868] NA NA NA NA ...
## $ articulated_parts  : chr [1:868] NA NA NA NA ...
## $ associated_parts   : chr [1:868] NA NA NA NA ...
## $ common_body_parts  : chr [1:868] NA NA NA NA ...
## $ rare_body_parts    : chr [1:868] NA NA NA NA ...
## $ feed_pred_traces   : chr [1:868] NA NA NA NA ...
## $ artifacts          : logi [1:868] NA NA NA NA NA NA ...
## $ component_comments : chr [1:868] NA NA NA NA ...
## $ pres_mode          : chr [1:868] NA NA NA NA ...
## $ preservation_quality : chr [1:868] NA NA NA NA ...
## $ spatial_resolution : chr [1:868] NA NA NA NA ...
## $ temporal_resolution : chr [1:868] NA NA NA NA ...
## $ lagerstätten       : chr [1:868] NA NA NA NA ...
## [list output truncated]
```

This looks reasonable. For example, we can see that our collection IDs are `numerical`, and our `identified_name` column contains `character` strings.

Now let's dive in further to look for inconsistencies in spelling, which could cause taxonomic names or geological units to be grouped separately when they are really the same thing. We'll start by checking for potential taxonomic misspellings.

We can use the `table()` function to look at the frequencies of various taxonomic names in the dataset. Here, inconsistencies like misspellings or antiquated taxonomic names might be recognised. We will check the columns `family`, `genus`, and `accepted_name`, the latter of which gives the name of the identification regardless of taxonomic level, and is the only column to give species binomials.

*# Tabulate the frequency of values in the "family" and "genus" columns*

```
table(fossils$family)
```

```
##
##      Alligatoridae      Crocodylidae      Gavialidae NO_FAMILY_SPECIFIED
##           205             61             34              230
##      Planocraniidae
##           24
```

```
table(fossils$genus)
```

```
##
##      Ahdeskatanka      Akanthosuchus      Aktiogavialis      Alligator
##           1             3             1              17
```

```
## Allognathosuchus      Argochampsia      Asiatosuchus      Asifcroco
##           129           3           32           1
## Astorgosuchus      Australosuchus      Baru      Borealosuchus
##           2           3           8           49
## Bottosaurus      Boverisuchus      Brachychampsia      Brachyuranochampsia
##           5           21           1           1
## Ceratosuchus      Chinaticampsus      Chrysochampsia      Crocodylus
##           5           1           1           42
## Diplocynodon      Dollosuchoides      Dongnanosuchus      Duerosuchus
##           65           1           1           1
## Dzungarisuchus      Eoalligator      Eocaiman      Eogavialis
##           1           3           5           2
## Eosuchus      Gavialosuchus      Gryposuchus      Hassiacosuchus
##           6           1           1           2
## Kambara      Kentisuchus      Krabisuchus      Leidyosuchus
##           4           3           3           1
## Lianghusuchus      Listrognathosuchus      Maomingosuchus      Maroccosuchus
##           2           1           4           3
## Megadontosuchus      Mekosuchus      Menatalligator      Mourasuchus
##           1           3           1           1
## Navajosuchus      Necrosuchus      Nihilichnus      Orientalosuchus
##           2           1           1           1
## Orthogenysuchus      Paratomistoma      Planocrania      Procaimanoidea
##           1           1           2           5
## Protoalligator      Protocaiman      Purussaurus      Qianshanosuchus
##           1           1           6           1
## Quinkana      Thecachampsia      Thoracosaurus      Tienosuchus
##           1           12           11           1
## Tomistoma      Trilophosuchus      Tsoabichi      Tzaganosuchus
##           4           1           2           1
## Ultrastenos      Wannaganosuchus
##           2           1
```

```
# Filter occurrences to those identified at species level, then tabulate species
# names
fossils_sp <- filter(fossils, accepted_rank == "species")
table(fossils_sp$accepted_name)
```

```
##
## Ahdeskatanka russlanddeutsche      Akanthosuchus langstoni
##           1           3
## Aktiogavialis puertoricensis      Alligator darwini
##           1           2
## Alligator gaudryi      Alligator hantoniensis
##           1           2
## Alligator prenasalis      Allognathosuchus heterodon
##           8           2
## Allognathosuchus mlynarskii      Allognathosuchus polyodon
##           1           2
## Allognathosuchus wartheni      Allognathosuchus woutersi
##           4           1
## Argochampsia krebsi      Asiatosuchus depressifrons
##           3           11
## Asiatosuchus germanicus      Asiatosuchus grangeri
##           3           1
```



##	Asiotosuchus nanlingensis	Asiotosuchus oenotriensis
##	4	1
##	Asifcroco retrai	Astorgosuchus bugtiensis
##	1	2
##	Australosuchus clarkae	Baru darrowi
##	3	2
##	Baru huberi	Baru wickeni
##	1	5
##	Borealosuchus acutidentatus	Borealosuchus formidabilis
##	2	17
##	Borealosuchus griffithi	Borealosuchus sternbergii
##	2	12
##	Borealosuchus wilsoni	Bottosaurus fustidens
##	2	2
##	Boverisuchus magnifrons	Boverisuchus vorax
##	2	17
##	Brachyuranochampsia eversolei	Ceratosuchus burdoshi
##	1	4
##	Chinatichampsus wilsonorum	Crocodylus ziphodon
##	1	2
##	Crocodylus aptus	Crocodylus elliotti
##	2	1
##	Crocodylus megarhinus	Diplocynodon darwini
##	3	1
##	Diplocynodon deponiae	Diplocynodon elavericus
##	3	1
##	Diplocynodon hantoniensis	Diplocynodon kochi
##	1	4
##	Diplocynodon levantinum	Diplocynodon muelleri
##	2	6
##	Diplocynodon ratelii	Diplocynodon remensis
##	2	2
##	Diplocynodon tormis	Dollosuchoides densmorei
##	4	1
##	Dongnanosuchus hsui	Duerosuchus piscator
##	1	1
##	Dzungarisuchus manacensis	Eoalligator chungyii
##	1	3
##	Eocaiman cavernensis	Eocaiman itaboraiensis
##	1	1
##	Eocaiman palaeocenicus	Eogavialis africanum
##	3	1
##	Eogavialis gavialoides	Eosuchus lerichei
##	1	1
##	Eosuchus minor	Gryposuchus jessei
##	5	1
##	Hassiacosuchus haupti	Kambara implexidens
##	1	1
##	Kambara molnari	Kambara murgonensis
##	1	1
##	Kambara taraina	Kentisuchus astrei
##	1	1
##	Kentisuchus spenceri	Krabisuchus siamogallicus
##	2	3

##	Lianghusuchus hengyangensis	Listrognathosuchus multidentatus
##	1	1
##	Maomingosuchus acutirostris	Maomingosuchus petrolica
##	1	2
##	Maroccosuchus zennaroi	Megadontosuchus arduini
##	3	1
##	Mekosuchus whitehunterensis	Menatalligator bergouniouxii
##	3	1
##	Navajosuchus mooki	Necrosuchus ionensis
##	2	1
##	Nihilichnus nihilicus	Orientalosuchus naduongensis
##	1	1
##	Orthogenysuchus olsenii	Paratomistoma courtii
##	1	1
##	Planocrania datangensis	Planocrania hengdongensis
##	1	1
##	Procaimanoidea kayi	Procaimanoidea utahensis
##	2	1
##	Protoalligator huiningensis	Protocaiman peligrensis
##	1	1
##	Qianshanosuchus youngi	Quinkana meboldi
##	1	1
##	Thecachampsia antiquus	Thecachampsia carolinensis
##	1	7
##	Thecachampsia marylandica	Thecachampsia sericodon
##	2	1
##	Thoracosaurus macrorhynchus	Thoracosaurus neocesariensis
##	1	5
##	Tienosuchus hsiangi	Tomistoma kerunense
##	1	1
##	Tomistoma tandoni	Tsoabichi greenriverensis
##	1	2
##	Tzaganosuchus infansii	Ultrastenos willisi
##	1	2
##	Wannaganosuchus brachymanus	
##	1	

Alternatively, we can use the `tax_check()` function in the `palaeoverse` package, which systematically searches for and flags potential spelling variation using a defined dissimilarity threshold.

```
# Check for close spellings in the "genus" column
tax_check(taxdf = fossils, name = "genus", dis = 0.1)
```

```
## $synonyms
## NULL
##
## $non_letter_name
## NULL
##
## $non_letter_group
## NULL
```

```
# Check for close spellings in the "accepted_name" column
tax_check(taxdf = fossils_sp, name = "accepted_name", dis = 0.1)
```

```
## $synonyms
```

```
## NULL
##
## $non_letter_name
## NULL
##
## $non_letter_group
## NULL
```

No names are flagged here, so that seems fine.

We can also check formatting and spelling using the `fossilbrush` package.

```
# Create a list of taxonomic ranks to check
fossil_ranks <- c("phylum", "class", "order", "family", "genus")
```

```
# Run checks
check_taxonomy(as.data.frame(fossils), ranks = fossil_ranks)
```

```
## Checking formatting [1/4]
## - formatting errors detected (see $formatting in output)
## Checking spelling [2/4]
## - no potential synonyms detected
## Checking ranks [3/4]
## - no cross-rank names detected
## Checking taxonomy [4/4]
## - conflicting classifications detected (see $duplicates in output)
## $formatting
## $formatting$'non-letter'
## $formatting$'non-letter'$phylum
## integer(0)
##
## $formatting$'non-letter'$class
## integer(0)
##
## $formatting$'non-letter'$order
## integer(0)
##
## $formatting$'non-letter'$family
## [1] 6 8 13 16 17 20 21 24 39 44 55 57 58 60 64 65 77 78
## [19] 82 83 85 86 89 90 265 269 291 292 294 296 297 298 299 300 301 302
## [37] 307 308 309 310 311 312 313 316 320 321 322 336 338 339 347 370 371 372
## [55] 373 374 375 379 392 393 394 395 399 408 409 414 421 429 432 433 438 439
## [73] 440 456 458 459 460 461 462 463 464 470 472 473 474 475 476 485 486 487
## [91] 488 489 492 493 494 499 500 501 502 503 504 505 506 507 508 509 510 511
## [109] 512 513 514 515 516 517 518 519 521 522 523 524 526 527 530 531 532 533
## [127] 534 535 543 544 545 546 547 548 549 552 556 557 563 564 565 571 572 574
## [145] 575 578 580 582 587 588 589 595 597 619 626 627 628 630 631 648 649 650
## [163] 651 652 656 657 660 667 669 670 671 673 674 675 676 677 686 692 697 701
## [181] 706 707 711 729 730 736 737 755 763 766 772 792 793 796 800 803 804 805
## [199] 806 808 809 811 817 818 819 820 821 822 823 824 825 826 827 828 842 845
## [217] 846 847 848 851 852 855 856 861 863 864 865 866 867 868
##
```

```

## $formatting$'non-letter'$genus
## integer(0)
##
##
## $formatting$'word-count'
## $formatting$'word-count'$phylum
## integer(0)
##
## $formatting$'word-count'$class
## integer(0)
##
## $formatting$'word-count'$order
## integer(0)
##
## $formatting$'word-count'$family
## integer(0)
##
## $formatting$'word-count'$genus
## integer(0)
##
##
##
## $ranks
## $ranks$crossed_adj
## $ranks$crossed_adj$'genus--family'
## character(0)
##
## $ranks$crossed_adj$'family--order'
## character(0)
##
## $ranks$crossed_adj$'order--class'
## character(0)
##
## $ranks$crossed_adj$'class--phylum'
## character(0)
##
##
## $ranks$crossed_all
## $ranks$crossed_all$genus
## character(0)
##
## $ranks$crossed_all$family
## character(0)
##
## $ranks$crossed_all$order
## character(0)
##
## $ranks$crossed_all$class
## character(0)
##
##
##
## $duplicates
## [1] taxon rank

```

```
## <0 rows> (or 0-length row.names)
```

As before, no major inconsistencies or potential spelling errors were flagged.

The PBDB has an integrated taxonomy system which limits the extent to which taxon name inconsistencies can arise. However, this is not the case for some other data fields. Let's check the formation names for our occurrences. This could be important if, for example, we wanted to include information about the number of formations from which our fossils are taken within the manuscript - a misspelling could inflate our count.

```
# Tabulating the frequency of each formation
table(fossils$formation)
```

```
##
##          ? Tepee Trail          "Brian Head"
##                1                1
##              Adaev              Akli
##                1                1
##            Aldearrubia          Aquia
##                1                7
##    Areniscas de Cabrerizos Argiles d'lignite du Soissonnais
##                1                4
##              Aridal            Arshanto
##                2                1
##            Ashawaq            Ashawq
##                1                1
##            Ashley            Avon Park
##                2                1
##            Aycross            Bandah
##                1                1
##            Barton Beds        Belgrade
##                1                1
##            Bellmunt            Birket Qarun
##                1                2
##            Black Peaks        Bohlen
##                5                1
##            Borgloon            Bridger
##                1                34
##            Bullion Creek      Calcaire de Saint-Ouen
##                2                1
##    Calcaries de Tàrrega        Canoe
##                3                1
##    Carl Creek Limestone        Chadron
##                14               11
##            Chambira            Chandler Bridge
##                7                4
##            Chitarwata          Clarno
##                7                2
##            Claron              Clayton
##                1                1
##    Clayton Limestone          Clinchfield
##                1                1
##    Cluj Limestone            Conglomérat de Cernay
##                2                1
##            Corçà              Corinda
##                1                1
##            Corrales           Couche I
```

##	7	1
##	Couche II	Creechbarrow Limestone
##	1	1
##	Cypress Hills	Daba'a
##	1	2
##	Dâncu	DeBeque
##	2	2
##	Denver	Devil's Graveyard
##	59	8
##	Dongjun	Drazinda
##	1	1
##	Elmore	Enspel
##	1	2
##	Entrala	Ergilin Dzo
##	6	1
##	Etadunna	Ferris
##	1	7
##	Fort Union	Gehannam
##	34	1
##	Golden Valley	Goler
##	3	4
##	Green River	Hannut
##	18	1
##	Harudi	Headon Beds
##	2	5
##	Headon Hill	Hetaoyuan
##	2	1
##	Heti	Honda
##	2	1
##	Hornerstown	Idam Unit
##	2	2
##	Irtys River	Itaboraí
##	1	1
##	Jebel Qatrani	Jewett
##	4	1
##	Kakara	Kamyshin
##	1	1
##	Karaginskaya	Kartal
##	1	1
##	Khirthar	Krabi Basin B2
##	1	5
##	Kressenberg	Kuldana
##	1	1
##	La Meseta	Laredo
##	2	3
##	Lignite-Sandstone	Limuping
##	2	1
##	Lingcha	Lizhuan
##	2	1
##	Lokone Sandstone	London Clay
##	3	1
##	Manasquan	Margaret
##	1	1
##	Maritsa	Mazaterón

##	2	1
##	Menat	Messel
##	1	8
##	Molasse de Carcassonne	Mortemer
##	2	1
##	Muñani	Na Duong
##	2	2
##	Nacimiento	Namba
##	10	5
##	Nanjemoy	Naran Bulak
##	4	1
##	Nongshan	North Horn
##	8	14
##	Ohio Creek	Orizzonte di Roncà
##	1	1
##	Panandhro	Paskapoo
##	1	2
##	Peñas Coloradas	Porcupine Hills
##	2	1
##	Pozo	Qasr el-Sagha
##	18	2
##	Quebrada de los Colorados	Ravenscrag
##	1	5
##	Rhems	Rundle
##	6	2
##	Sables de Bracheux	Sables du Castrais
##	2	2
##	Salamanca	San Sebastian
##	7	2
##	San Sebastián	Santo Spirito
##	2	1
##	Sarmiento	Scollard
##	1	1
##	Selsey	Shark River
##	1	1
##	Silveirinha	Slim Buttes
##	1	1
##	Solimões	Stone City
##	8	2
##	Subathu	Sylhet Limestone
##	4	1
##	Syzran	Tallahatta
##	2	1
##	Teberemt	Tienen
##	1	8
##	Tremembé	Tulloch
##	1	3
##	Turbuța	Twiggs
##	1	1
##	Uinta	Upper Headon Beds
##	5	1
##	Upper Variegated Sequence	Valea Nadășului
##	1	1
##	Vaugirard	Viens Marl

```
##          4          1
##      Vincentown      Viștea Limestone
##          9          3
##      Wanghudun      Wasatch
##          2          17
##      Washakie      Williamsburg
##          5          2
##      Willwood      Wind River
##      168          1
##      Xinyu      Yazoo
##          1          1
##      Yegua      Yongning
##          1          2
##      Youganwo
##          2
```

```
# Checking for spelling variation
tax_check(taxdf = fossils, name = "formation", dis = 0.1)
```

```
## $synonyms
##   group      greater      lesser count_greater count_lesser
## 1    A      Ashawq      Ashawaq          1          1
## 2    C      Clarno      Claron          2          1
## 3    C      Couche II      Couche I          1          1
## 4    S San Sebastián San Sebastian          2          2
##
## $non_letter_name
## [1] "Devil's Graveyard"      "Calcaire de Saint-Ouen"
## [3] "\"Brian Head\""          "Argiles d'lignite du Soissonnais"
## [5] "Qasr el-Sagha"          "Daba'a"
## [7] "Krabi Basin B2"          "? Tepee Trail"
## [9] "Lignite-Sandstone"
##
## $non_letter_group
## [1] "\" " "?"
```

Here, we see four flagged formation name pairs, which vary in how problematic they seem. For example “Couche II” and “Couche I” make sense as an adjacent pair of formations, and are likely not a mistake. However, we also see “San Sebastián” and “San Sebastian”, which should be unified under the same name, either with or without the accent included.

We can also look for inconsistencies in the interval names and dates. For example, an interval could be spelled two different ways, or have multiple associated dates.

```
# Again, we can use the table() function to tabulate the frequency of each
# interval
table(fossils$early_interval)
```

```
##
##      Arikareean      Arshantan      Bartonian      Bridgerian      Cf1
##          1          1          33          58          1
##      Cf2          Cf3      Chadronian      Chattian      Clarkforkian
##          2          2          12          21          28
##      Danian      Duchesnean      Early Eocene Early Oligocene Early Paleocene
##          32          2          14          23          2
##      Early Uintan      Eocene      Graybullian      Late Eocene      Late Oligocene
```



```
##           9           10           1           23           37
## Late Paleocene Late Uintan Lostcabinian Lutetian Middle Eocene
##           9           6           3           53           17
##           MP 10           MP 11           MP 13           MP 17A           MP 18
##           3           8           1           2           2
##           MP 21           MP 22           MP 23           MP 24           MP 25
##           1           1           1           1           1
##           MP 26           MP 28           MP 30           Oligocene           Paleogene
##           2           2           3           8           1
##           Priabonian           Pu1           Puercan           Rupelian           Selandian
##           28           1           94           16           4
##           Thanetian           Ti3           Ti5           Tiffanian           Torrejonian
##           21           3           1           18           16
##           Uintan           Wasatchian           Whitneyan           Ypresian
##           5           185           1           38
```

```
table(fossils$late_interval)
```

```
##
##           Bartonian           Bridgerian           Chattian           Clarkforkian           Duchesnean
##           4           2           9           1           1
##           Early Eocene Early Oligocene           Late Eocene           Lutetian           Middle Eocene
##           2           1           7           7           1
##           MN 1           MP 17B           MP 19           Pliocene           Priabonian
##           1           2           2           8           5
##           Rupelian           Selandian           Thanetian           Torrejonian           Ypresian
##           4           2           2           1           3
```

```
# What are the ages associated with each interval? Any issues with the date
# formats?
```

```
fossils %>%
```

```
  group_by(early_interval) %>% # Grouping by interval
```

```
  distinct(max_ma) %>% # Pulling out the distinct values
```

```
  arrange(early_interval) # Putting list in order to scan for inconsistencies
```

```
## # A tibble: 54 x 2
```

```
## # Groups:   early_interval [54]
```

```
##   early_interval max_ma
```

```
##   <chr>          <dbl>
```

```
## 1 Arikareean    29.5
```

```
## 2 Arshantan     56
```

```
## 3 Bartonian     41.2
```

```
## 4 Bridgerian    50.5
```

```
## 5 Cf1           56.2
```

```
## 6 Cf2           56.2
```

```
## 7 Cf3           56.2
```

```
## 8 Chadronian    37
```

```
## 9 Chattian      27.8
```

```
## 10 Clarkforkian 56.2
```

```
## # i 44 more rows
```

```
fossils %>%
```

```
  group_by(late_interval) %>%
```

```
  distinct(min_ma) %>%
```

```
  arrange(late_interval)
```

```
## # A tibble: 39 x 2
## # Groups:   late_interval [21]
##   late_interval  min_ma
##   <chr>         <dbl>
## 1 Bartonian      37.7
## 2 Bridgerian     46.2
## 3 Chattian       23.0
## 4 Clarkforkian   54.9
## 5 Duchesnean     37
## 6 Early Eocene   47.8
## 7 Early Oligocene 27.8
## 8 Late Eocene    33.9
## 9 Lutetian       41.2
## 10 MN 1         22.4
## # i 29 more rows
```

Finally, let's check the coordinates using the `CoordinateCleaner` package.

```
# Is the coordinate formatting OK? Numeric and part of a lat/long system?
coord_check <- cc_val(fossils, lat = "lat", lon = "lng")
```

```
## Testing coordinate validity
```

```
## Removed 0 records.
```

```
paleocoord_check <- cc_val(fossils, lat = "paleolat", lon = "paleolng")
```

```
## Testing coordinate validity
```

```
## Removed 0 records.
```

```
# Testing for coordinates with equal lat and long
```

```
coord_check <- cc_equ(coord_check, lat = "lat", lon = "lng")
```

```
## Testing equal lat/lon
```

```
## Removed 0 records.
```

```
paleocoord_check <- cc_equ(paleocoord_check, lat = "paleolat", lon = "paleolng")
```

```
## Testing equal lat/lon
```

```
## Removed 0 records.
```

```
# Near any centroids of modern-day political units or biodiversity centers?
```

```
coord_cent <- cc_cen(coord_check, lat = "lat", lon = "lng", value = "flagged")
```

```
## Testing country centroids
```

```
## Flagged 0 records.
```

```
coord_cent <- cc_inst(coord_check, lat = "lat", lon = "lng", value = "flagged")
```

```
## Testing biodiversity institutions
```

```
## Flagged 0 records.
```

```
# Any zeros?
```

```
coord_check <- cc_zero(coord_check, lat = "lat", lon = "lng")
```

```
## Testing zero coordinates
```

```
## Removed 0 records.
```

```
# We can also test for temporal mismatches
```

```
coord_check <- coord_check[!is.na(coord_check$min_ma),]
```

```
coord_check <- coord_check[!is.na(coord_check$max_ma),]
coord_check <- cf_equal(coord_check, min_age = "min_ma", max_age = "max_ma")
```

```
## Testing age validity
```

```
## Removed 0 records.
```

Overall, our dataset looks to be in good shape!

## Rule 8: Identify and handle duplicates

Our next step is to remove duplicates. This is an important step for count data, as duplicated values will artificially inflate our counts. Here, the function `dplyr::distinct()` is incredibly useful, as we can provide it with the columns we want it to check, and it removes rows for which data within those columns is identical.

First, we will remove *absolute* duplicates: by this, we mean occurrences within a single collection which have identical taxonomic names. This can occur when, for example, two species are named within a collection, one of which is later synonymised with the other.

```
# Show number of rows in dataset before duplicates are removed
nrow(fossils)
```

```
## [1] 868
```

```
# Remove occurrences with the same collection number and `accepted_name`
fossils <- distinct(fossils, collection_no, accepted_name, .keep_all = TRUE)
```

```
# Show number of rows in dataset after duplicates are removed
nrow(fossils)
```

```
## [1] 844
```

The number of rows dropped from 868 to 844 - this means 24 of our occurrences were absolute duplicates, which have now been removed.

Next, we can look at geographic duplicates. We mentioned earlier that sometimes PBDB collections are entered separately for different beds from the same locality, and this means that the number of collections can be higher than the number of geographic sampling localities. Let's check whether this is the case in our dataset.

```
# Remove duplicates based on geographic coordinates
fossils_localities <- distinct(fossils, lng, lat, .keep_all = TRUE)
```

```
# Compare length of vector of unique collection numbers with and without this
# filter
length(unique(fossils$collection_no))
```

```
## [1] 703
```

```
length(unique(fossils_localities$collection_no))
```

```
## [1] 518
```

Here we can see that our original dataset contains 703 collections, but once we remove latitude-longitude repeats, this drops to 518. This means that we have 518 geographically distinct localities, across which 703 fossil sampling events have taken place.

If we are interested in taxonomic diversity, we can also look at repeated names in our dataset. For example, we might want to identify taxa which are represented multiple times in order to then return to the literature and check that they definitely represent the same taxon. We can do this by flagging species names which are represented more than once in the dataset.

```
# Update dataset of occurrences identified to species level
fossils_sp <- filter(fossils, accepted_rank == "species")
```

```
# Identify and flag taxonomic duplicates
fossils_sp <- fossils_sp %>%
  group_by(accepted_name) %>%
  mutate(duplicate_flag = n() > 1)
```

```
# Show counts of flagged occurrences
table(fossils_sp$duplicate_flag)
```

```
##
## FALSE  TRUE
##    57   185
```

Our counts show 57 FALSE values, indicating that 57 species are represented by a single occurrence. We also have 185 TRUE values, for which the species is represented two or more times. We can then filter our dataset to those flagged, and sort them by their name, enabling easier checking.

```
# Filter table to flagged occurrences
fossils_sp <- filter(fossils_sp, duplicate_flag == TRUE)

# Sort table by genus name
fossils_sp <- arrange(fossils_sp, accepted_name)
```

We recommend iteratively enacting Rules 4 to 7, because if data are altered or filtered at any point, this can change the overall summary statistics, and affect how we perceive the data. Once you are happy with your dataset, it is time to perform any analytical processes relating to your research question - in our simple example, this is quantifying the latitudinal distribution of our occurrences over time.

```
# List geological stages in dataset
time_bins <- unique(fossils$interval_name)

# Create object to hold values
lat_ranges <- c()

# For each stage, identify maximum and minimum latitudinal values
for (i in 1:length(time_bins)) {
  one_stage <- filter(fossils, interval_name == time_bins[i])
  max <- max(one_stage$paleolat)
  min <- min(one_stage$paleolat)
  midpoint <- one_stage$bin_midpoint[1]
  lat_ranges <- rbind(lat_ranges, c(time_bins[i], midpoint, min, max))
}


# Clean up dataframe
lat_ranges <- as.data.frame(lat_ranges)
colnames(lat_ranges) <- c("stage", "midpoint", "min", "max")
lat_ranges$midpoint <- as.numeric(lat_ranges$midpoint)
lat_ranges$max <- as.numeric(lat_ranges$max)
lat_ranges$min <- as.numeric(lat_ranges$min)

# Plot the palaeolatitudinal range by stage
ggplot(lat_ranges, aes(x = midpoint, ymin = min, ymax = max)) +
  geom_ribbon() +
  labs(x = "Age (Ma)",
```

```

y = "Palaeolatitude (°)" +
scale_x_reverse() +
geom_hline(yintercept = 0) +
coord_geo(dat = "stages", expand = TRUE)

```



## Rule 9: Report your data and cleaning efforts

It is important to report the steps you took in cleaning and processing your data. When any code used to achieve this is presented alongside the manuscript, this is particularly easy, as any specific details which readers may be interested in can be checked within the code. As a result, the written description can be kept brief. An example outlining our workflow above is as follows:

“We downloaded all fossil occurrence data for the taxon ‘Crocodylia’, dated to the time interval ‘Paleogene’ (66–23 Ma) from the Paleobiology Database on 21st February 2025 (see supplementary data). Taxonomic names and geographic coordinates were checked for inconsistencies and potential errors, and any duplicates were removed (see supplementary code). After cleaning, the dataset included 844 occurrences, spread across 518 discrete geographic sampling locations.”

We can also create a reference list containing all of the sources for the occurrences in our dataset.

```

# Generate list of reference numbers
refs <- sort(unique(fossils$reference_no))
# Reformat the list
refs <- paste("ref:", refs, ",", sep = "")
refs <- paste(refs, collapse = "")
refs <- gsub('.{1}$', '', refs)
# Create API call
ref_url <- paste("https://paleobiodb.org/data1.2/refs/list.csv?ref_id=",
  refs, sep = "")

```

```

# Use API call (this is not enacted here)
ref_table <- read_csv(ref_url)
# Save table of references for supplement

```

```
write_csv(ref_table, "data/data_references.csv")
```

Depending on the preferences of our journal, we can then add these to our reference list, or provide the table as part of the paper's data supplement.

## Rule 10: Deposit your data and workflow

Now that we've completed our data cleaning and exploration, we want to make sure that our data and workflow are well-documented and easily accessible to others. To this end, we have developed this vignette, which we have been maintaining within a [repository on GitHub](#) throughout the development of our project. This vignette contains all of the code used to clean and explore our data, as well as explanations of each step. We have also included the [raw data file](#) in the repository, so that others can download and use it for their own analyses. We have made sure to include all of the necessary metadata in the raw data file, so that others can understand where the data came from and how it was processed. With all of this, other researchers are able to run this vignette locally and follow our proposed rules step-by-step.

To ensure our workflow is citable, we have linked our GitHub repository to [Zenodo](#), which now archives each [release](#) of our repository and provides a DOI for citation. We have also included a [copyleft license](#) to ensure that others can use and build upon our work.

## Summary

To conclude, we hope that this vignette illustrates how the ten rules can be put into action. The example given here is simple, and we would expect that more thorough exploration and cleaning would be necessary for a scientific project, ideally with these steps being tailored to the dataset and research question involved.